

- (b) State in English the complement of the problem HAM-PATH.
 - (c) Prove or disprove: DNF-SAT is the complement of CNF-SAT (See Exercise 6.7).
 - (d) A Boolean formula ϕ is called a *tautology* if it evaluates to *true* for all possible truth assignments of its variables. Prove that deciding whether a Boolean formula is a tautology is in coNP.
 - (e) A Boolean formula ϕ is called a *contradiction* if it evaluates to *false* for all possible truth assignments of its variables. Prove that deciding whether a Boolean formula is a contradiction is in coNP. What is its complement problem called?
- 6.33 Prove that $P \subseteq NP \cap \text{coNP}$. (**Remark:** It is not known whether the containment is proper.)
- 6.34 Let SMALL-FACTOR denote the problem of deciding whether a given positive integer n has a factor no larger than a given positive integer k . Prove that SMALL-FACTOR is in $NP \cap \text{coNP}$.
- 6.35 Prove that if SMALL-FACTOR can be solved in polynomial time, integers can be factored in polynomial time too, and conversely.
- 6.36 Prove that if $NP \neq \text{coNP}$, then $P \neq NP$. (**Remark:** It is widely believed that $NP \neq \text{coNP}$.)
- 6.37 Prove that if an NP-Complete problem is in coNP, then $NP = \text{coNP}$.
- 6.38 Define coNP-Hard and coNP-Complete problems as those to which there exist polynomial-time reductions from *all* problems in coNP. Prove that a problem is coNP-Complete if and only if its complement is NP-Complete. State a few coNP-Complete problems.
- 6.39 Prove that the TAUTOLOGY problem (see Exercise 6.32(d)) is coNP-Complete.
- 6.40 Prove that the CONTRADICTION problem (see Exercise 6.32(e)) is coNP-Complete.
- 6.41 Prove or disprove: The halting problem is coNP-hard.
- 6.42 Let FSAT denote the problem of finding a satisfying assignment of a Boolean formula ϕ or return failure if ϕ is unsatisfiable. Prove that FSAT can be solved in polynomial-time if and only if $\text{SAT} \in P$.

Approximation algorithms

- 6.43 Consider the optimization version of the set covering problem of Exercise 6.13. That is, given a finite set S and a collection of k subsets S_1, S_2, \dots, S_k of S , we intend to find out a cover of S (from the given collection) of size as small as possible. Let us denote this optimization problem by MIN-SET-COVER.
- Let $S = \{x_1, x_2, \dots, x_n\}$ be of size n , and let f_i be the count of the subsets S_j containing the element x_i . Finally, let $f = \max(f_1, f_2, \dots, f_n)$. (The counts f_i are the frequencies of the elements, and f is the maximum frequency.)
- Design a polynomial-time f -approximation algorithm for MIN-SET-COVER. Establish that your algorithm achieves an approximation ratio of f . Determine whether this approximation ratio is tight.
- 6.44 Prof. Myopia proposes the following approximation algorithm for solving the MAX-CUT problem (Exercise 6.27).
1. Start with an arbitrary partition S, T of V .

2. Repeat the following two steps until no further vertex movement is possible:
 - (a) For each vertex $v \in S$, check whether the cut $(S - v, T + v)$ has more cross edges than (S, T) ; and if so, delete v from S and include v in T .
 - (b) For each vertex $v \in T$, check whether the cut $(S + v, T - v)$ has more cross edges than (S, T) ; and if so, delete v from T and include v in S .
 3. Return S, T .
- (a) Prove that Prof. Myopia's algorithm runs in polynomial time (in the input size).
 - (b) Prove or disprove: Prof. Myopia's algorithm outputs the optimal solution for bipartite graphs.
 - (c) Prove that the approximation ratio of Prof. Myopia's algorithm is $1/2$.
 - (d) Demonstrate that this approximation ratio is tight (suggest an *infinite* family of graphs).
- 6.45** [Weighted Max-Cut Problem] Let $G = (V, E)$ be an undirected graph with each edge $e \in E$ having a (positive) cost c_e . The cost of a cut (S, T) (where S and T are disjoint subsets of V , and $S \cup T = V$) is $c(S, T) = \sum_{\substack{u \in S \\ v \in T}} c_{(u,v)}$. The task is to produce a cut (S, T) of G that maximizes $c(S, T)$. Adapt Prof. Myopia's algorithm (Exercise 6.44) to a $\frac{1}{2}$ -approximation algorithm for the weighted max-cut problem.
- 6.46** Let $G = (V, E)$ be a connected undirected graph. You make a DFS traversal of G starting from any arbitrary vertex. Let T be the DFS tree produced by the traversal. Take C to be the set of all internal (that is, non-leaf) nodes in T .
- (a) Prove that C is a vertex cover for G .
 - (b) Prove that the determination of C using this method is a 2-approximation algorithm for the MIN-VERTEX-COVER problem.
- 6.47** Adapt the approximation algorithm of Exercise 6.46 to a general undirected graph (not necessarily connected).
- 6.48** Consider the following algorithm for EUCLIDEAN TSP, that iteratively builds a closed tour C for the traveling salesperson. First, we choose two cities u, v such that $d(u, v)$ is smallest among all pairs of (different) cities. We start with $C = (u, v)$ (that is, go from u to v , and then back to u). Next, we run a loop, each iteration of which adds a new city to the currently constructed tour C . This city is chosen as follows. Find $i \in C$ and $j \notin C$ such that $d(i, j)$ is minimized. Let k be the city following i in the current tour C . Replace i, k by i, j, k in C . Repeat until C contains all the cities. Prove that this is a 2-approximation algorithm.
- 6.49** Consider the knapsack problem introduced in Section 25.1.4. First, sort the objects in decreasing order of p_i/w_i . Call this sorted list O_1, O_2, \dots, O_n . Let k be the index such that $\sum_{i=1}^k w_i \leq C$ and $\sum_{i=1}^{k+1} w_i > C$. We have seen that the greedy algorithm that outputs $\{1, 2, \dots, k\}$ can be arbitrarily bad. Now, let j be the index such that p_j is maximum. Augment the greedy algorithm to output $\{1, 2, \dots, k\}$ or $\{j\}$, whichever gives better profit. Prove that this is a $1/2$ -approximation algorithm for the knapsack problem.
- 6.50** Let us use the notations of Exercise 6.49. Modify the augmented greedy algorithm to output $\{1, 2, \dots, k\}$ or $\{k+1\}$, whichever gives better profit. Prove that this is again a $1/2$ -approximation algorithm for the knapsack problem.
- 6.51** Consider the balanced set partition problem introduced in Exercise 6.18.
- (a) Prove that this problem cannot have any polynomial-time approximation algorithm (unless $P = NP$).

- (b) We reformulate the problem as follows. Let A be the set with the larger sum. We want to minimize $\sum_{a \in A} a$ subject to the constraint that $\sum_{a \in A} a \geq \frac{1}{2} \sum_{x \in S} x$. Argue that there exists a trivial 2-approximation algorithm for this optimization problem.
- 6.52** Let $G = (V, E)$ be an undirected graph with each vertex v having a positive weight $w(v)$. The *weighted vertex cover problem* deals with the computation of a vertex cover C of G such that $\sum_{v \in C} w(v)$ is as small as possible.
- (a) Prove that the 2-approximation algorithm of Section 25.1.1 fails for the weighted vertex cover problem, that is, given any positive constant Δ , there exists a graph for which the approximation factor is $> \Delta$.
- (b) Prove that the 2-approximation algorithm of Exercise 6.47 fails too for the weighted vertex cover problem.
- 6.53** Assume that $P \neq NP$. Prove that there cannot exist any ρ -approximation algorithm for the BIN-PACKING problem (Exercise 6.22) for $\rho < 3/2$.
- 6.54** The *next-fit* strategy for BIN-PACKING keeps on adding objects to the most recently opened bin so long as possible. When further placements are not possible, the last bin is closed, a new bin is opened, and the next object is put in the newly opened bin.
- (a) Prove that the next-fit strategy is a 2-approximation algorithm.
- (b) Prove that the approximation ratio 2 is tight, that is, given any $\varepsilon > 0$, there exists an input instance for which the algorithm achieves an approximation ratio $> 2 - \varepsilon$.
- 6.55** The *first-fit* strategy for BIN-PACKING keeps all bins open. The objects are inserted in the sequence they appear in S . Each item is attempted to be inserted in the earliest open bin which can accommodate the object. If no opened bin can accommodate the object, a new bin is opened and the object is put in this newly opened bin. Prove that the first-fit strategy is a 2-approximation algorithm.
- [H] **6.56** The *first-fit decreasing* strategy for BIN-PACKING first sorts the objects in the decreasing order of their weights, and inserts them in that sorted sequence following the first-fit strategy. Prove that this is a $3/2$ -approximation algorithm.
- 6.57** The *linear bin-packing problem* deals with inserting the objects in the order they appear in the input. That is, for $i < j$, the j -th object cannot be packed in a bin opened earlier than the bin containing the i -th object. Argue that the next-fit strategy solves the linear bin-packing problem exactly.
- [H] **6.58** Let $G = (V, E)$ be the complete undirected graph such that each pair (u, v) of vertices is associated with a distance $d(u, v) \geq 0$. Assume that $d(u, u) = 0$ for all u , $d(u, v) = d(v, u)$ for all u, v , and the distances satisfy the triangle inequality: $d(u, w) \leq d(u, v) + d(v, w)$ for all u, v, w . Let $S \subseteq V$ with $|S| = k$. For each $u \in V$, define $d(u, S) = \min_{v \in S} d(u, v)$. The *k-center problem* deals with the determination of a set S of size k such that $\max_{u \in V} d(u, S)$ is minimized. This problem is used for clustering the points in V with S playing the role of the set of cluster centers. Propose a 2-approximation algorithm for the k -center problem.
- 6.59** [Makespan scheduling] There are m identical machines, and n jobs with running times t_1, t_2, \dots, t_n assigned to the machines. The machines start operation at time $t = 0$, and carry out the jobs assigned to them in a non-preemptive fashion without any waiting time between two consecutive jobs. The time when all the jobs finish is called the *makespan*. The task is to distribute the jobs to the machines in such a way that the makespan is minimized.
- First sort t_1, t_2, \dots, t_n in non-increasing order, and assume that $t_1 \geq t_2 \geq \dots \geq t_n \geq 0$. Now, for $i = 1, 2, \dots, n$, assign the i -th job to the machine which finishes at the earliest (based upon the

assignment of the first $i - 1$ jobs). This is called *greedy makespan scheduling*. Deduce that this is a 2-approximation algorithm.

- 6.60** Suppose that t_1, t_2, \dots, t_n are not sorted as in the greedy makespan scheduling algorithm of Exercise 6.59. The rest of the algorithm remains the same. Prove that this is still a 2-approximation algorithm.
- 6.61** An algorithm is called *pseudo-polynomial-time*, if its running time is a polynomial in the size of the *unary* representation of the input. We call an NP-Complete problem *weakly NP-Complete* if it admits a pseudo-polynomial-time algorithm. Prove that the subset-sum problem (Exercise 6.14) is weakly NP-Complete.
- 6.62** Prove that the partition problem (Exercise 6.17) is weakly NP-Complete.
- 6.63** Prove that the balanced set partitioning problem (to be more precise, its decision version, see Exercise 6.18) is weakly NP-Complete.
- 6.64** Prove that the SUBSET-DIFF problem (Exercise 6.19) is weakly NP-Complete.
- 6.65** Prove that the 2-SET-SUM problem (Exercise 6.20) is weakly NP-Complete.
- 6.66** Prove that the 2-SET-PARTITION problem (Exercise 6.21) is weakly NP-Complete.
- 6.67** Prove that the ferry-loading problem (Exercise 6.24) is weakly NP-Complete.
- [H³] **6.68** An *asymptotic polynomial time approximation scheme* (APTAS) is a ρ -approximation algorithm such that $\rho \rightarrow 1 + \varepsilon$ as $\text{OPT} \rightarrow \infty$ for some constant $\varepsilon > 0$. The best approximation ratio that we can achieve for BIN-PACKING is $3/2$ (Exercises 6.53 and 6.56). However, BIN-PACKING has an APTAS. More precisely, design an algorithm that given a constant $\varepsilon \in (0, 1/2]$ produces an output $m \leq (1 + \varepsilon)\text{OPT} + 1$.

Randomized algorithms

- 6.69** Recall that a Las Vegas algorithm always outputs correct answers. Let us investigate a similar class of algorithms which may sometimes report *failure*. But whenever the algorithms succeed, the answer output is correct. Let us call such an algorithm a Las Vegas' algorithm. Let A' be a Las Vegas' algorithm for some problem. Using this algorithm, design a Las Vegas algorithm A that never outputs *failure*. Assume that A' outputs *failure* with probability $\leq 1/2$. Express the expected running time of A in terms of the expected running time of A' .
- 6.70** Suppose that a one-sided error Monte Carlo algorithm A has error probability $\leq 1/2$. Let $\delta > 0$ be a small real number. How many times you should run A in order to reduce the error probability to $< \delta$?
- 6.71** Consider the problem of finding the i -th smallest element in an array A of n integers. Ms. Lucky proposes the following algorithm to solve this problem. She chooses a (uniformly) random element x of A . She then uses the partitioning algorithm of Quick Sort on A with respect to the pivot x . Suppose that x is placed in the k -th position after the partitioning (counting starts from 1). If $k = i$, the algorithm returns x . If $k > i$, then a recursive call is made on the smaller subarray (of size $k - 1$) and with the same i . Finally, if $k < i$, then a recursive call is made on the larger subarray (of size $n - k$) with i replaced by $i - k$. Deduce that the expected running time of Ms. Lucky's algorithm is $O(n \log n)$. (Notice that this running time may depend upon i (in addition to n). In your calculations, you may suitably ignore this dependence.) What is the worst-case running time of Ms. Lucky's algorithm?